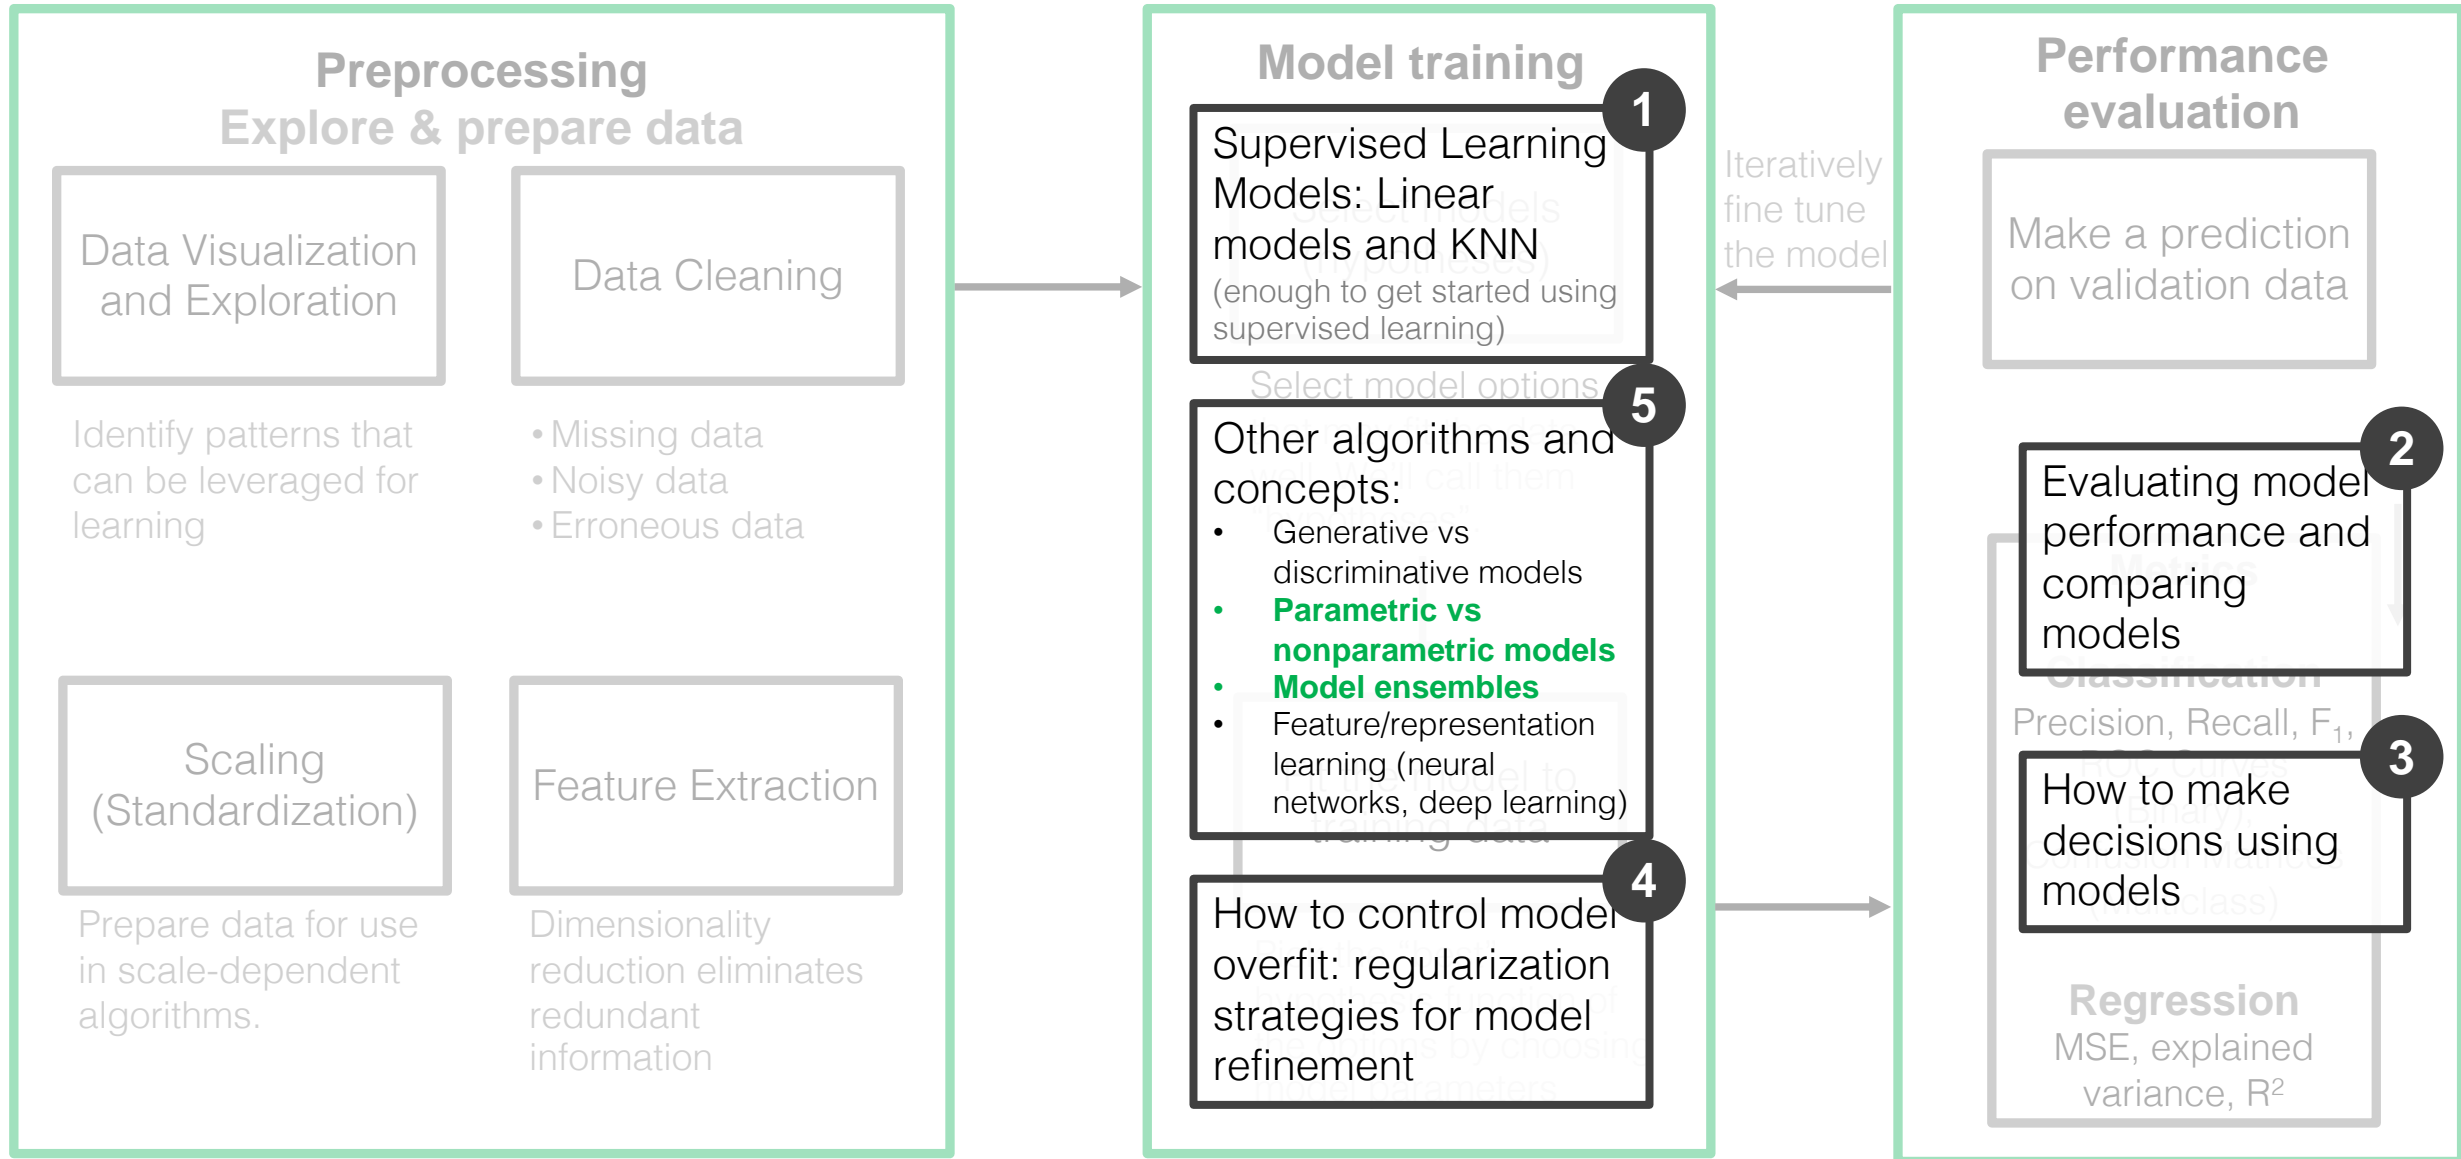


Tree-based Models and Ensembles

Supervised learning in practice



Supervised Learning Techniques

Covered so far

Linear Regression

K-Nearest Neighbors

Perceptron

Logistic Regression

Linear Discriminant Analysis

Quadratic Discriminant Analysis

Naïve Bayes

Decision Trees and Random Forests

Ensemble methods (bagging, boosting, stacking)

Parametric vs Nonparametric techniques

Non-parametric Models

Complexity of the model grows with the size of the training data

- K-Nearest Neighbors
- Decision Trees

Parametric Models


Fixed number of parameters (i.e. a fixed structure)

- Linear regression
- Logistic regression
- LDA, QDA
- Naïve Bayes with Gaussian likelihoods

Classification and Regression Trees (CART)

Classification trees = decision trees

Predicting promotions of
salaried employees




Classification and Regression Trees (CART)

Predicting promotions of salaried employees

1

Find the best “split” in any one feature (that best classifies the data) that divides the region in two

● Promoted
● Not promoted




Classification and Regression Trees (CART)

Predicting promotions of salaried employees

1

Find the best “split” in any one feature (that best classifies the data) that divides the region in two



Classification and Regression Trees (CART)

Predicting promotions of salaried employees

1


Find the best “split” in any one feature (that best classifies the data) that divides the region in two

2

Continue splitting regions (1 feature at a time) until a stopping criterion is reached (e.g. there are at most N samples in any region)


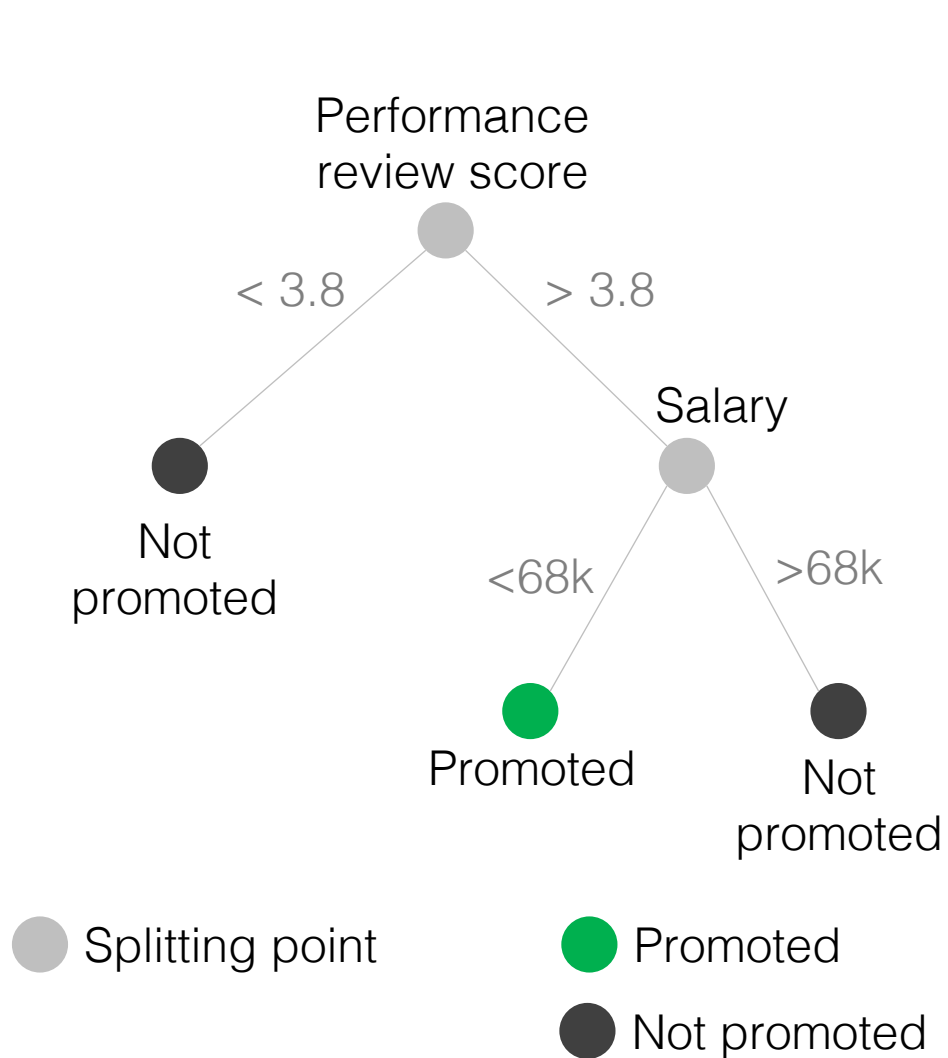
**Greedy, recursive
binary tree**

● Promoted
● Not promoted



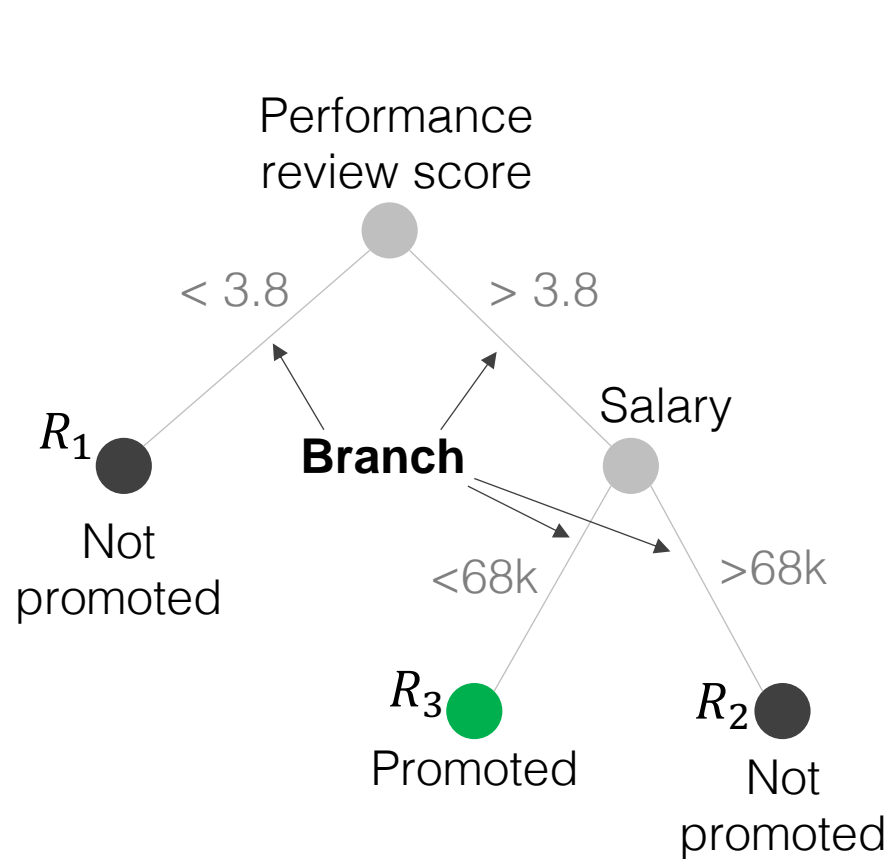
Classification and Regression Trees (CART)

Tree representation:



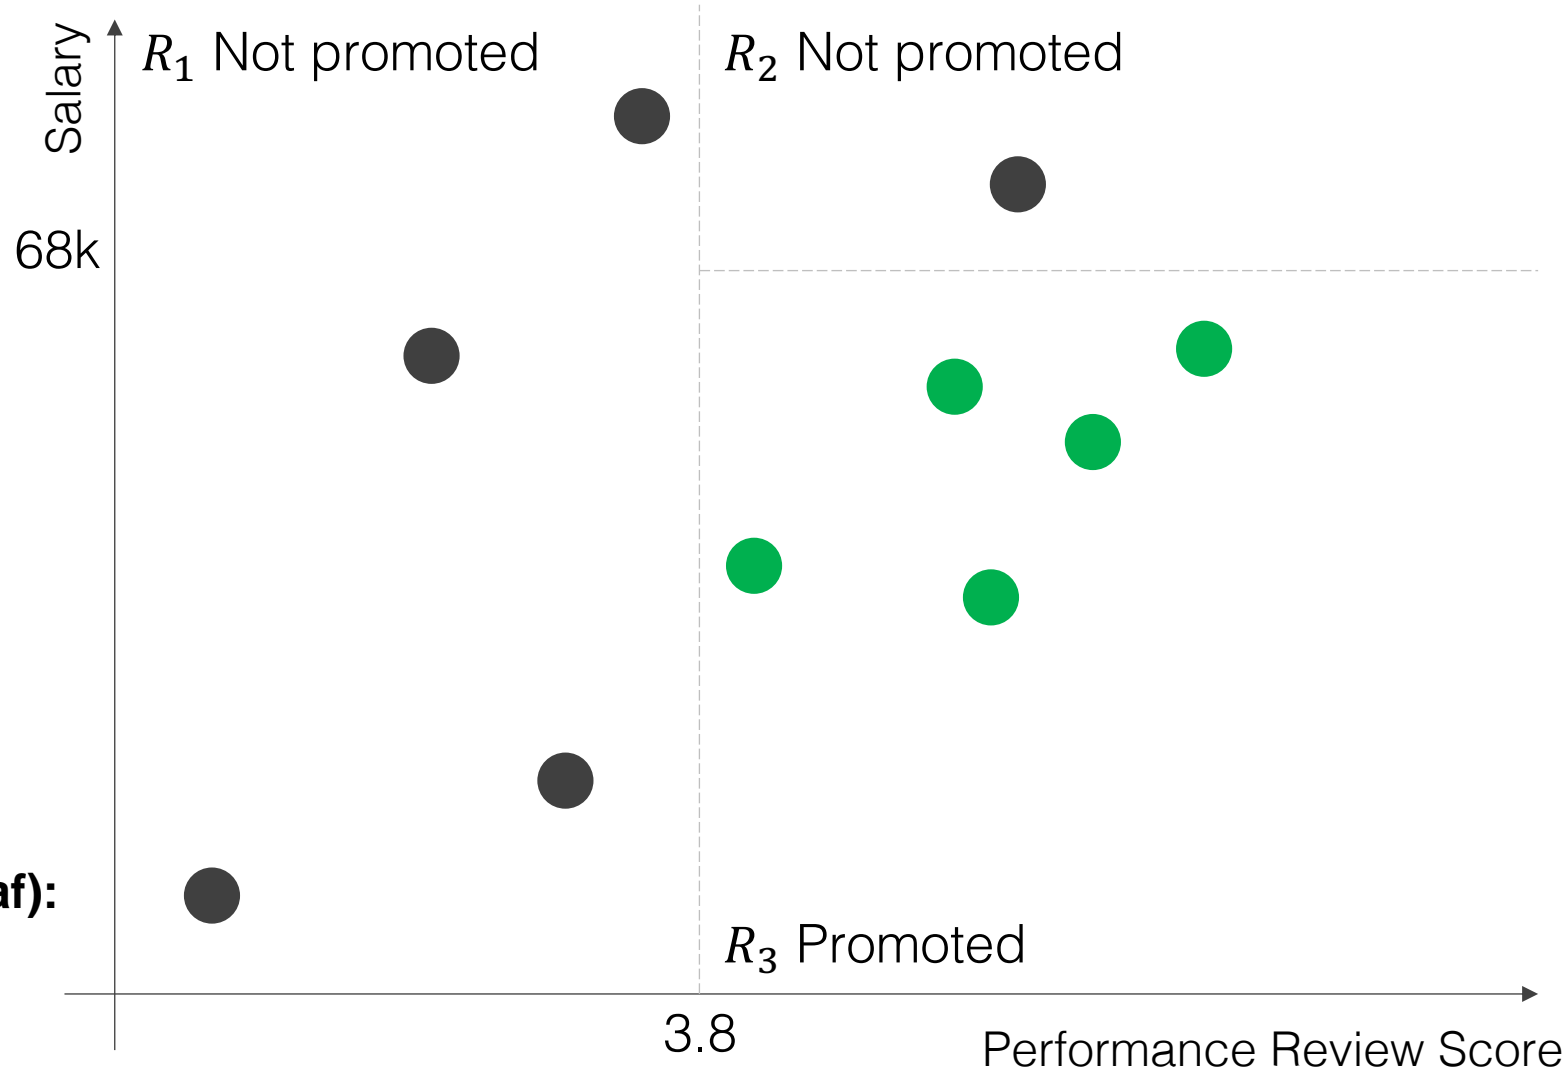
Classification and Regression Trees (CART)

Tree representation:



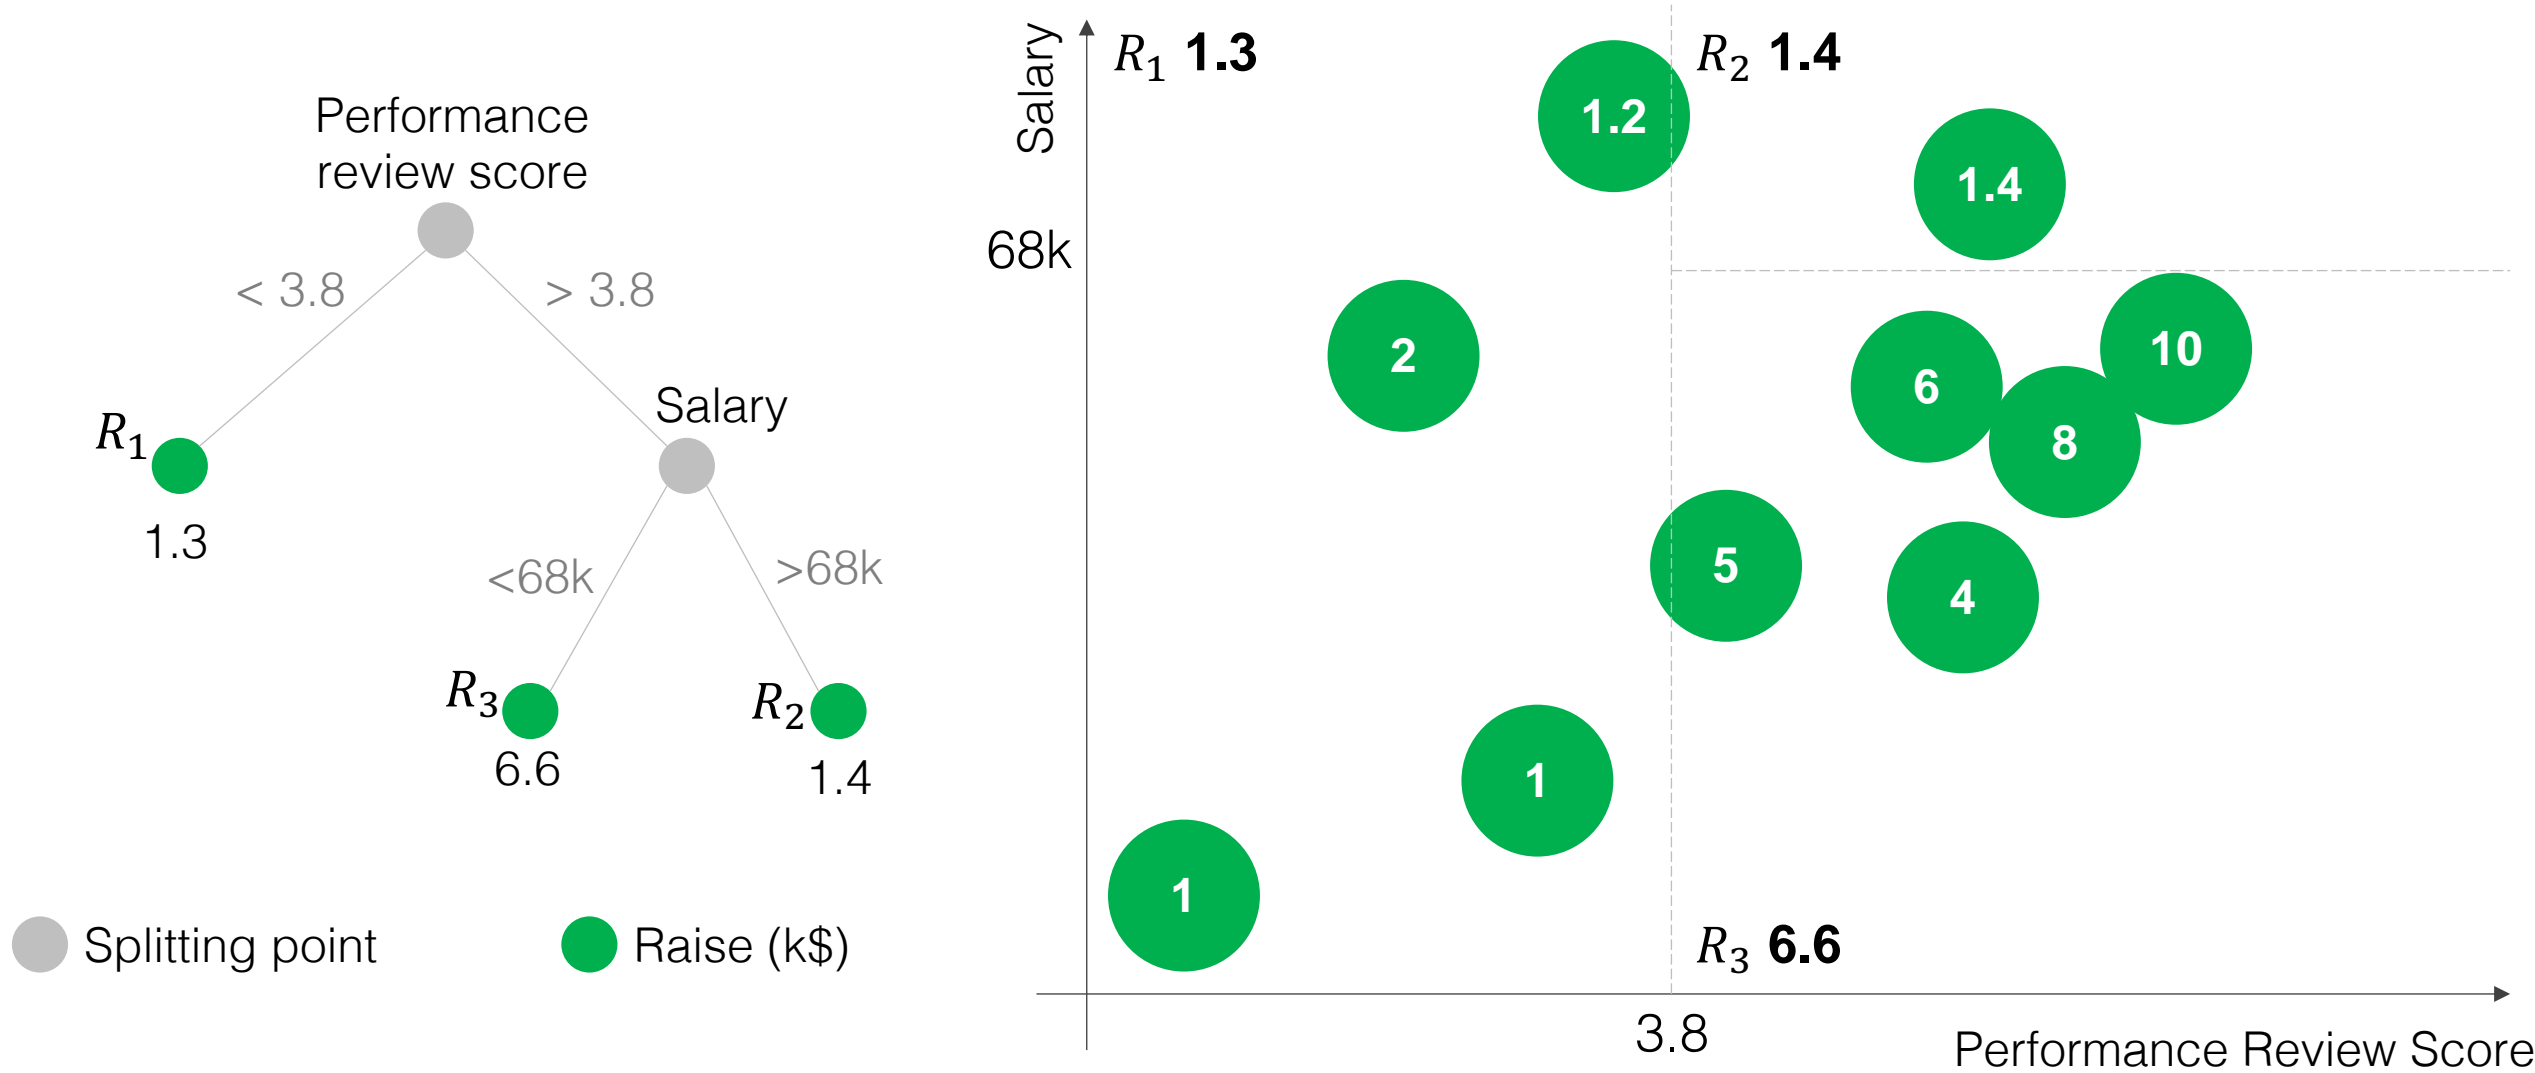
Internal node:
● Splitting point

Terminal node (leaf):
● Promoted
● Not promoted



The Regression Setting

In this case, each region is represented by an average of the values it contains



How do we determine which split to make?

Pick the split that reduces the error/cost criterion most after the split

Splitting criterion

$$C = \sum_{r=1}^{R_{tot}} Q(r)$$

Regression

Mean square error

$$Q_{MSE}(r) = \sum_{i \in R_r} (y_i - \hat{y}_{R_r})^2$$

y_i = training data response i
 \hat{y}_{R_r} = mean value in region r ,
(where R_r is the set of samples in region r)

Classification

Misclassification rate

$$Q_{Misclass}(r) = 1 - \max_k (\hat{p}_{rk})$$

Gini impurity

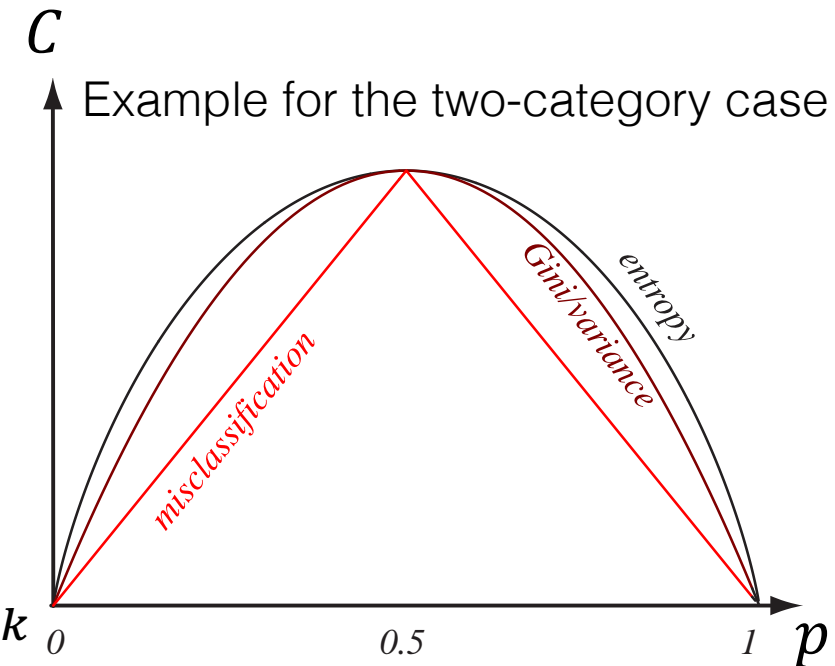
Measure of how often a randomly chosen element from the set would be incorrectly labeled if it was randomly labeled

$$Q_{Gini}(r) = \sum_{k=1}^K \hat{p}_{rk}(1 - \hat{p}_{rk})$$

Cross-entropy

$$Q_{entropy}(r) = - \sum_{k=1}^K \hat{p}_{rk} \log \hat{p}_{rk}$$

\hat{p}_{rk} = proportion of training observations in the r^{th} region from the k^{th} class



Duda, Hart, and Stork., Pattern Classification

How to measure quality of split for classification?

\hat{p}_{rk} = proportion of training observations in the r^{th} region from the k^{th} class

Class 1 ●
Class 2 ●

For each region:

Misclassification rate

$$Q_{\text{Misclass}}(r) = 1 - \max_k (\hat{p}_{rk})$$

A

0.333

B

0.167

Gini impurity

$$Q_{\text{Gini}}(r) = \sum_{k=1}^K \hat{p}_{rk}(1 - \hat{p}_{rk})$$

0.444


0.278

Cross-entropy

$$Q_{\text{entropy}}(r) = - \sum_{k=1}^K \hat{p}_{rk} \log \hat{p}_{rk}$$

0.637

0.450



Tree Pruning

Trees have the tendency to overfit the data

Consider the stopping rule: stop splitting once there is only 1 observation in each region (leads to complete overfit)


Pruning the tree back reduces this overfit (removing splits after the tree is formed)

Pruning can be optimized through a penalty on the number of terminal nodes:

$$C_{Prune} = \sum_{j=1}^T \sum_{i \in R_j} (y_i - \hat{y}_{R_j})^2 + \alpha T$$


penalty on number of terminal nodes number of terminal nodes

Decision Tree




Pruning example

Pruned Tree




Original Tree

Example: heart disease classification







Performance



James et al., An Introduction to Statistical Learning

Linear model



Classification Tree

Struggle when the boundary is not parallel to an axis

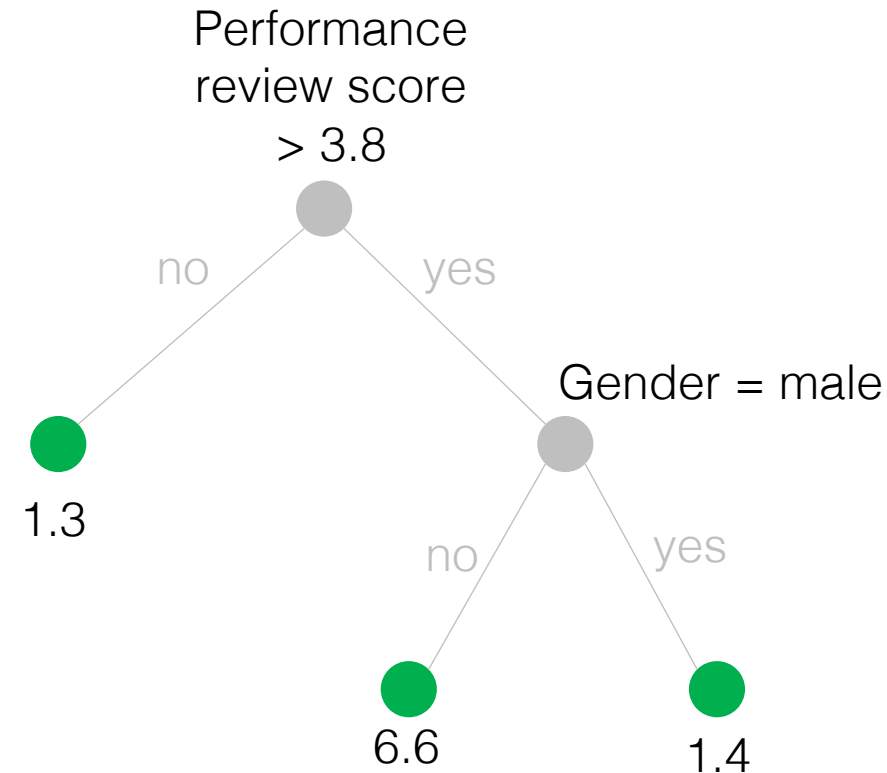
...nonlinear feature transforms could help...

James et al., An Introduction to Statistical Learning

Pros/Cons

Numerical data

Categorical data



Pros:

Trees easily handle multiple types of data

Trees are easy to interpret

Cons:

Trees do not typically have the same level of predictive accuracy of many methods

Tend to overfit
(have high variance)

Ensemble learning

Combining models to improve performance beyond any individual model alone

Bagging (bootstrap aggregation)

Random forests (tree-specific modification of bagging)

Gradient boosting

Reducing Variance or Bias through ensembles

Bagging

Models in ensemble:

high variance, low bias
(i.e. overfit models)

Effect of aggregating:

Reduce variance through averaging output

Boosting

high bias, low variance
(i.e. underfit models, “weak learners”)


Reduce bias through sequentially fitting models to previous model errors

Bagging

Bootstrap aggregation

Trees **overfit** (have high variance). Averaging over observations **reduces variance**

Recall bootstrap sampling (sampling with replacement):



Bagging

Decision trees are a popular base model

Bootstrap aggregation

- 1 Create a random bootstrap sample from the training data
- 2 Train a model on that bootstrap sample and call it $\hat{f}_i(\mathbf{x})$
- 3 Repeat 1 and 2 until we have B models trained on different bootstrap samples
- 4 Take the average of the output for our new model estimate:

$$\hat{f}_{bag}(\mathbf{x}) = \frac{1}{B} \sum_{i=1}^B \hat{f}_i(\mathbf{x})$$

(for classification models we can get the average class confidence or take a majority vote)

Bagging

Tree Number:

1

2

3

4

Observations
Included:
(out of 1-9)

[1,2,3,3,8]

[1,2,4,7,7]

[1,5,6,8,9]

[2,2,2,4,9]

Features list:

[A, B, C, D]





[A, B, C, D]

[A, B, C, D]

[A, B, C, D]

First split:

Trees:



Variable Importance

Decision trees are very interpretable, but this is lost with bagging


We can construct another measure called “variable importance” to **compare feature contributions**

1


Calculate the total amount the error (or impurity) decreased by splitting on each feature.

2

Average over all the trees resulting from bagging



Random Forests



Random Forests

A **small tweak on bagging**

Random forests
decorrelate
the bagged trees

Decision trees are constructed greedily

This can lead to highly correlated trees

“Strong” features will typically be split before moderately strong predictors.

Each time a split is considered, a **random subset of m features** is selected as candidates from the full set of p features

Typically chose: $m = \sqrt{p}$

(If $m = p$, then we would be back to the bagging approach)


Bagging

Observations
Included:
(out of 1-9)

Features list:

$[1, 2, 3, 3, 8]$


$[A, B, C, D]$



Random forests

$[1, 2, 3, 3, 8]$

$[A, B, C, D]$







Boosting

Decision trees are a popular base model

Bagging created trees that were designed to be as independent as possible

Boosting involves building trees **sequentially**, each building on the errors of the last



We build consecutive models, each fit to the residuals of the last model

We sum models output to get the boosted prediction

$$\hat{f}_{boost}(x) = \hat{f}_1(x) + \hat{f}_2(x) + \hat{f}_3(x)$$

Boosting

Boosting for regression trees

- 1 Select the number of models to train, B , and learning rate λ
- 2 Set $\hat{f}(\mathbf{x}) = 0$ and $r_i = y_i$ for all the training data
- 3 Fit a tree, $\hat{f}_b(\mathbf{x})$ to the residuals, r_i (with d splits)
- 4 Update $\hat{f}(\mathbf{x}) \leftarrow \hat{f}(\mathbf{x}) + \lambda \hat{f}_b(\mathbf{x})$
- 5 Update the residuals $r_i \leftarrow r_i - \lambda \hat{f}_b(\mathbf{x}_i)$
- 6 Output the boosted model:
$$\hat{f}(\mathbf{x}) = \sum_{b=1}^B \lambda \hat{f}_b(\mathbf{x})$$

λ slows down the learning process to avoid overfitting

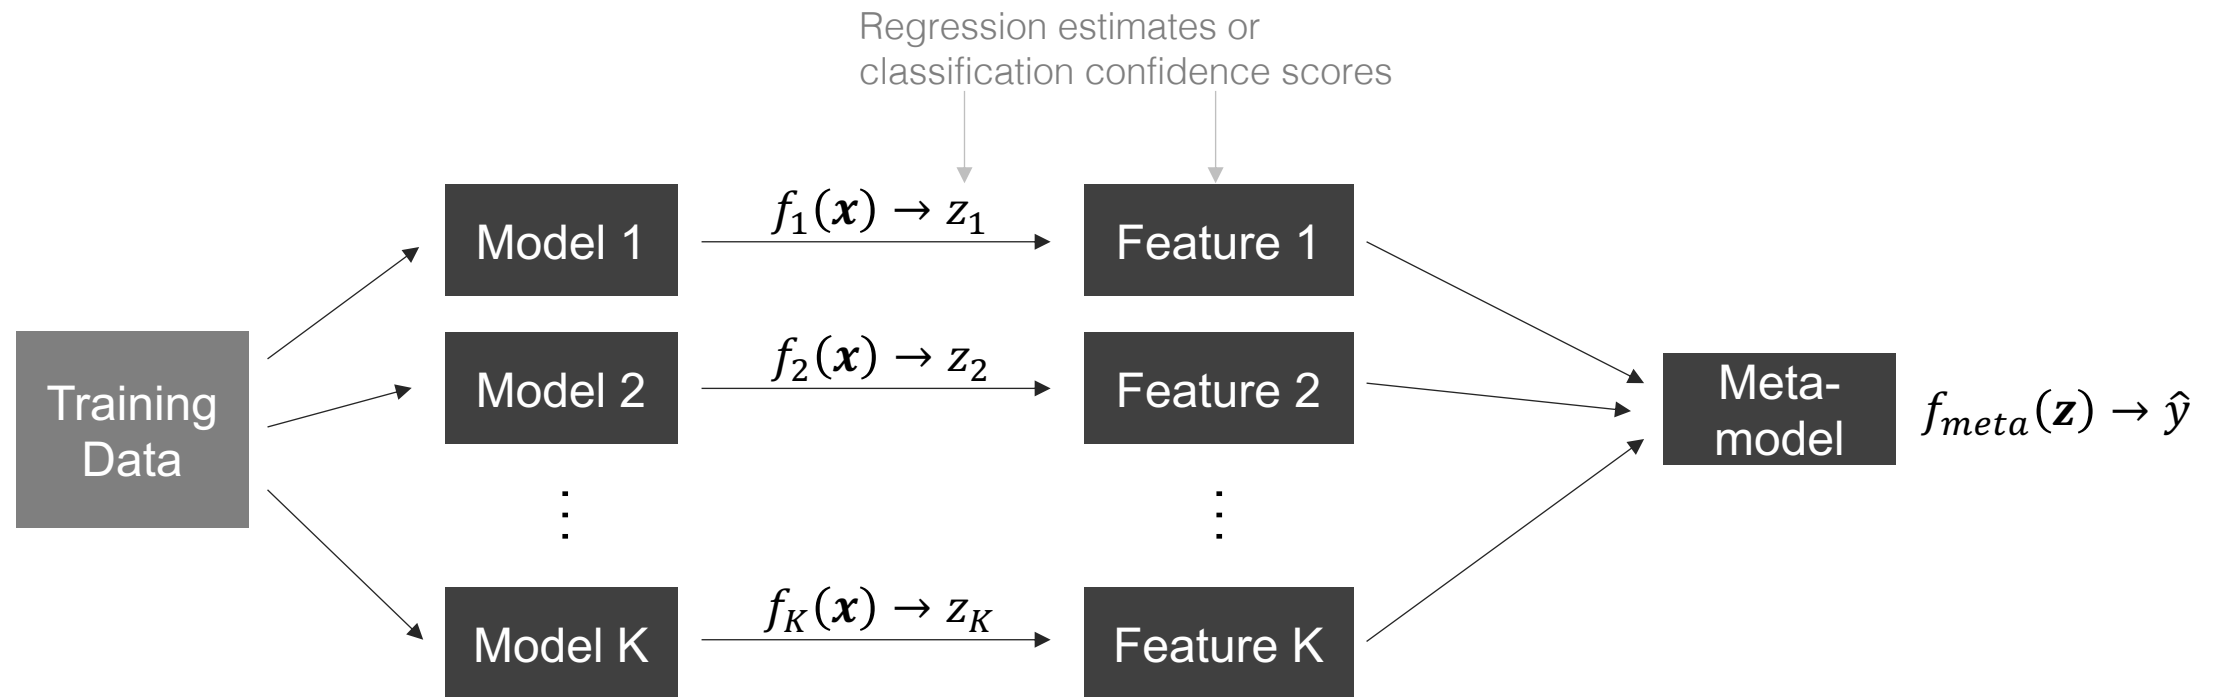
Often this is just a small number of splits (a stump)

Repeat B times

Model Stacking

Train multiple supervised learning techniques (could be different models)

THEN Train a supervised learning technique that includes the **outputs** of the other models as **features**



Supervised Learning Techniques

Covered so far

- Linear Regression
- K-Nearest Neighbors
- Perceptron
- Logistic Regression
- Linear Discriminant Analysis
- Quadratic Discriminant Analysis
- Naïve Bayes
- Decision Trees and Random Forests
- Ensemble methods (bagging, boosting, stacking)

Can be used with numerous machine learning techniques, often CART

Appropriate for:
● Classification
● Regression