

Deep Learning

Deep learning

Representation learning with a hierarchy of concepts

Those concepts are represented by layers in a neural network model

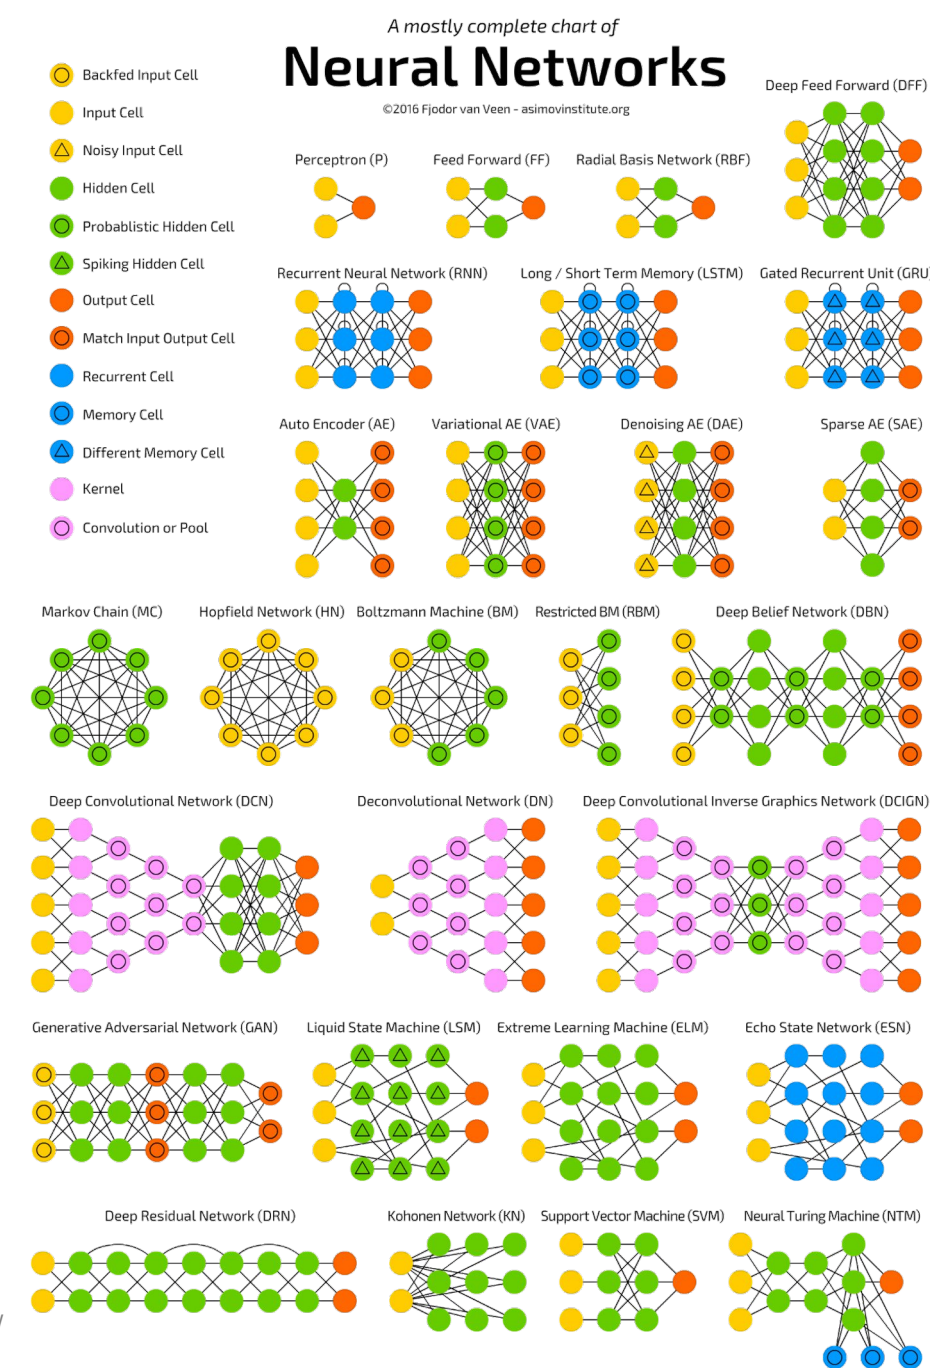
Types of Deep Learning Tools

Convolutional Neural Networks

Autoencoders

Recurrent Neural Networks (e.g. LSTMs)



Generative Adversarial Networks (GANs)



Azimov Institute: <http://www.asimovinstitute.org/neural-network-zoo/>


Convolutional Neural Networks

AlexNet



Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E. Hinton. "Imagenet classification with deep convolutional neural networks." Advances in neural information processing systems. 2012.

Key



Convolutional Neural Networks





Image from the Mathworks

Data: \mathbf{x}

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

$*$

Weights: \mathbf{w}

1	1	1
0	0	0
-1	-1	-1

$=$

Output: $\mathbf{x} * \mathbf{w}$

2D Convolution

Data: x

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

Weights: w

1	1	1
0	0	0
-1	-1	-1

*

=

Output: $x * w$

Computing
one output
value:

$$\begin{aligned}
 &1 \cdot 1 + 1 \cdot 2 + 1 \cdot 5 + \\
 &0 \cdot 0 + 0 \cdot 2 + 0 \cdot 3 + \\
 &(-1) \cdot 4 + (-1) \cdot 5 + (-1) \cdot 5
 \end{aligned}$$

2D Convolution

Data: x

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

Weights: w

1	1	1
0	0	0
-1	-1	-1

*

=

Output: $x * w$

-6			

Computing
one output
value:

$$\begin{aligned}
 &1 \cdot 1 + 1 \cdot 2 + 1 \cdot 5 + \\
 &0 \cdot 0 + 0 \cdot 2 + 0 \cdot 3 + \\
 &(-1) \cdot 4 + (-1) \cdot 5 + (-1) \cdot 5 = -6
 \end{aligned}$$

2D Convolution

Data: X

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

Weights: w

1	1	1
0	0	0
-1	-1	-1

*

=

Output: $X * w$

-6	-11		

Computing
one output
value:

$$\begin{aligned}
 &1 \cdot 2 + 1 \cdot 5 + 1 \cdot 1 + \\
 &0 \cdot 2 + 0 \cdot 3 + 0 \cdot 2 +
 \end{aligned}$$

$$(-1) \cdot 5 + (-1) \cdot 5 + (-1) \cdot 9 = -11$$

2D Convolution

Data: X

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

Weights: w

1	1	1
0	0	0
-1	-1	-1

$*$

$=$

Output: $X * w$

-6	-11	-12	

Computing
one output
value:

$$\begin{aligned}
 &1 \cdot 5 & + & 1 \cdot 1 & + & 1 \cdot 4 & + \\
 &0 \cdot 3 & + & 0 \cdot 2 & + & 0 \cdot 0 & +
 \end{aligned}$$

$$(-1) \cdot 5 + (-1) \cdot 9 + (-1) \cdot 8 = -12$$

2D Convolution

Data: X

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

Weights: w

1	1	1
0	0	0
-1	-1	-1

*

=

Output: $X * w$

-6	-11	-12	-11

Computing
one output
value:

$$\begin{aligned}
 &1 \cdot 1 + 1 \cdot 4 + 1 \cdot 2 + \\
 &0 \cdot 2 + 0 \cdot 0 + 0 \cdot 0 +
 \end{aligned}$$

$$(-1) \cdot 9 + (-1) \cdot 8 + (-1) \cdot 1 = -11$$

2D Convolution

Data: X

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

Weights: w

1	1	1
0	0	0
-1	-1	-1

*

=

Output: $X * w$

-6	-11	-12	-11
-7			

Computing
one output
value:

$$\begin{aligned}
 &1 \cdot 0 & + & 1 \cdot 2 & + & 1 \cdot 3 & + \\
 &0 \cdot 4 & + & 0 \cdot 5 & + & 0 \cdot 5 & +
 \end{aligned}$$

$$(-1) \cdot 6 + (-1) \cdot 3 + (-1) \cdot 4 = -7$$

2D Convolution

Data: X

1	2	5	1	4	2
0	2	3	2	0	0
4	5	5	9	8	1
6	3	4	2	3	1
0	1	9	8	7	2
2	3	5	5	5	6

6 x 6

Weights: w

1	1	1
0	0	0
-1	-1	-1

3 x 3

*

=

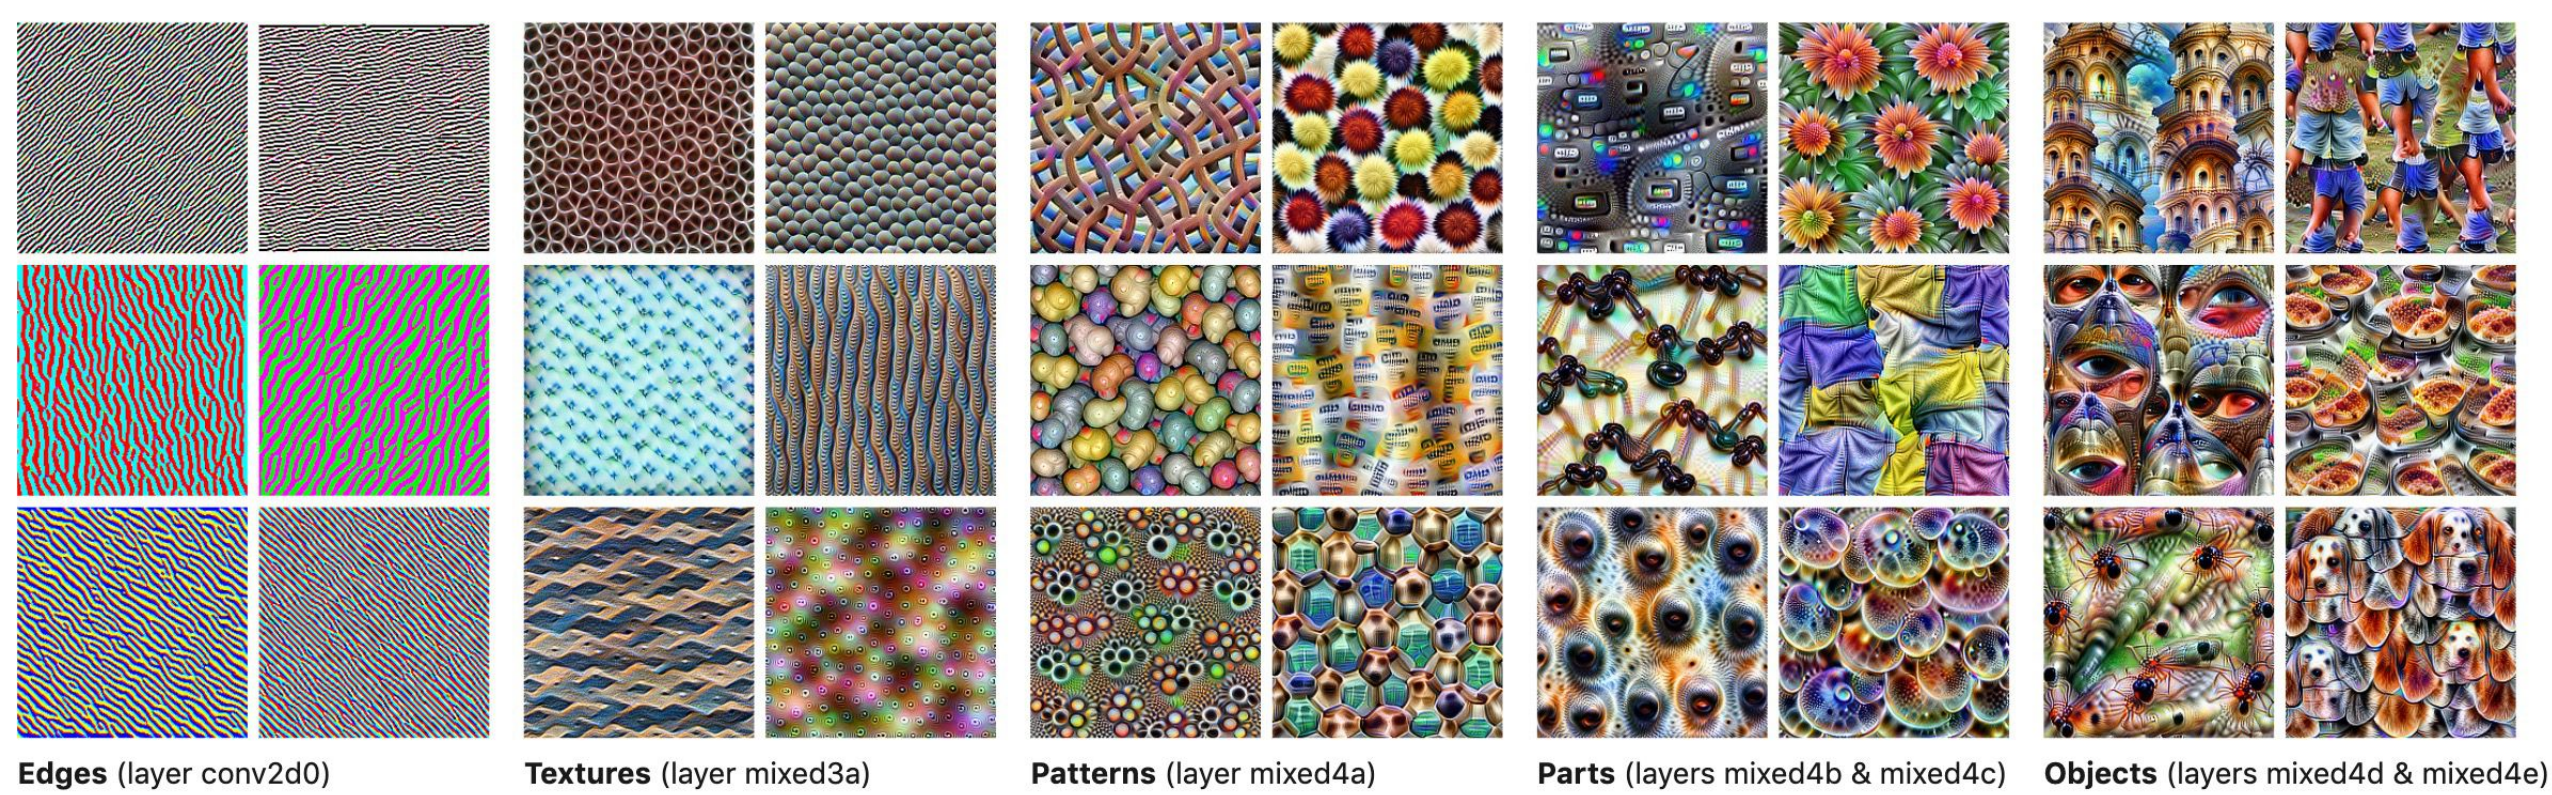

Output: $X * w$

-6	-11	-12	-11
-7	-2	-2	-4
4	1	-2	1
3	-4	-6	-10

4 x 4

2D Convolution

What features do layers respond to?




Olah et al, 2017: <https://distill.pub/2017/feature-visualization/>

Features


Dataset Examples show us what neurons respond to in practice



Optimization isolates the causes of behavior from mere correlations. A neuron may not be detecting what you initially thought.



Baseball—or stripes?
mixed4a, Unit 6



Animal faces—or snouts?
mixed4a, Unit 240



Clouds—or fluffiness?
mixed4a, Unit 453



Buildings—or sky?
mixed4a, Unit 492

Olah et al, 2017: <https://distill.pub/2017/feature-visualization/>


Resources on Visualization of Features

Feature visualization: <https://distill.pub/2017/feature-visualization/>

Building blocks of interpretability: <https://distill.pub/2018/building-blocks/>


Activation Activation Atlases: <https://distill.pub/2019/activation-atlas/>

Convolution Layer



From Fei-Fei Li, Justin Johnson, and Serena Young. CS231n, 2018


For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:



We stack these up to get a “new image” of size 28x28x6!

From Fei-Fei Li, Justin Johnson, and Serena Young. CS231n, 2018

Preview: ConvNet is a sequence of Convolution Layers, interspersed with activation functions




$$\text{Parameters} = (5 \times 5 \times 3) \times 6 = 450$$

$$(5 \times 5 \times 6) \times 10 = 1,500$$


From Fei-Fei Li, Justin Johnson, and Serena Young. CS231n, 2018

1 x 1 Convolution Explained




1x1 CONV
with 32 filters

(each filter has size
1x1x64, and performs a
64-dimensional dot
product)




From Fei-Fei Li, Justin Johnson, and Serena Young. CS231n, 2018

Max Pooling




max pool with 2x2 filters
and stride 2




From Fei-Fei Li, Justin Johnson, and Serena Young. CS231n, 2018

Fully Connected Layer



AlexNet




Input
11x11 conv, 96
5x5 conv, 256
max pool
3x3 conv, 384
max pool
3x3 conv, 384
3x3 conv, 256
max pool
FC 4096
FC 4096
FC 1000
softmax

Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E. Hinton. "Imagenet classification with deep convolutional neural networks." Advances in neural information processing systems. 2012.

Key

Input or output layer
Convolutional Layer
Fully Connected Layer
max pooling layer


AlexNet
(2012)



VGG16
(2014)



VGG19
(2014)



Note: an activation function is applied to the output of each layer

Fewer layers, larger filters

Key

Input or output layer

Convolutional Layer

Fully Connected Layer

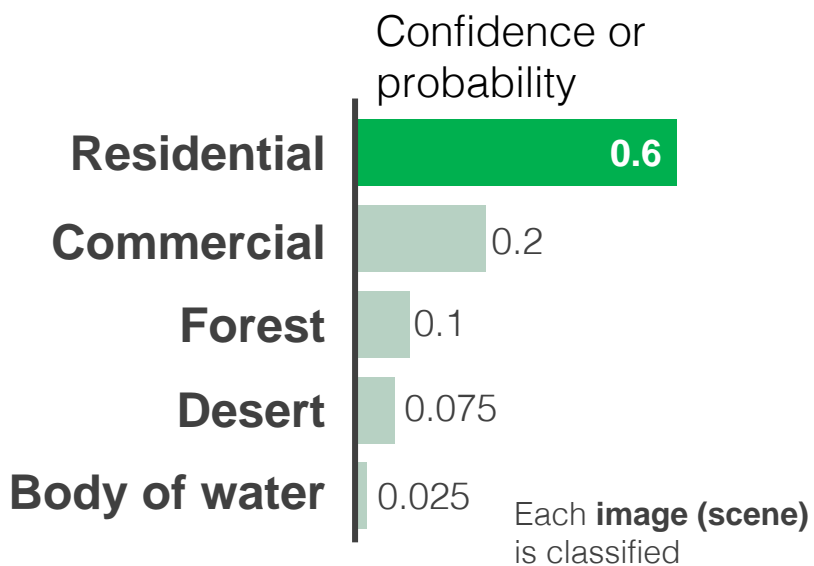
max pooling layer

CNN Architectures

Adapted from Fei-Fei Li, Justin Johnson, and Serena Young. CS231n, 2018



Scene
classification



Object
detection

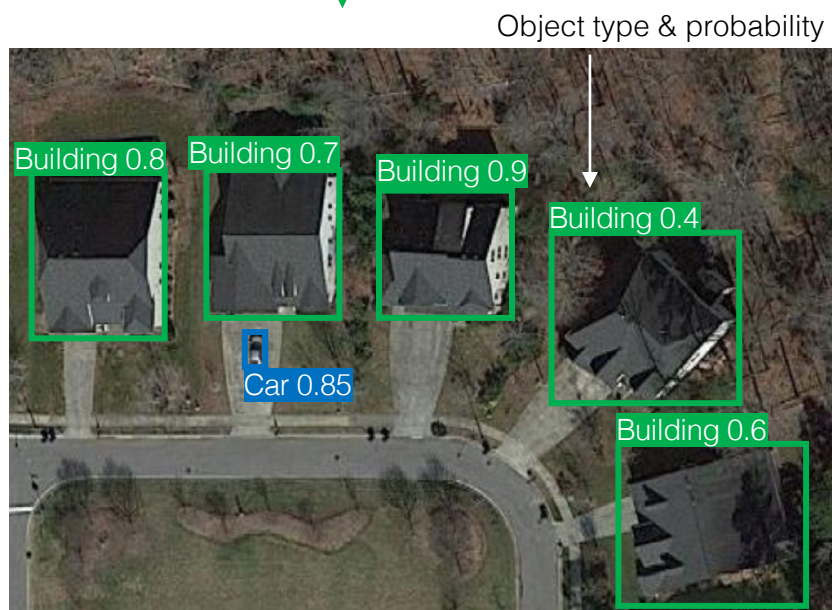


Image
segmentation






Scene classification



AlexNet
VGG
GoogLeNet
ResNet

Inception
DenseNet
SqueezeNet
EfficientNet

Confidence or probability



Object detection



Faster/Fast/R-CNN
Mask R-CNN
YOLO
Single Shot Detector (SSD)
RetinaNet

Object type & probability

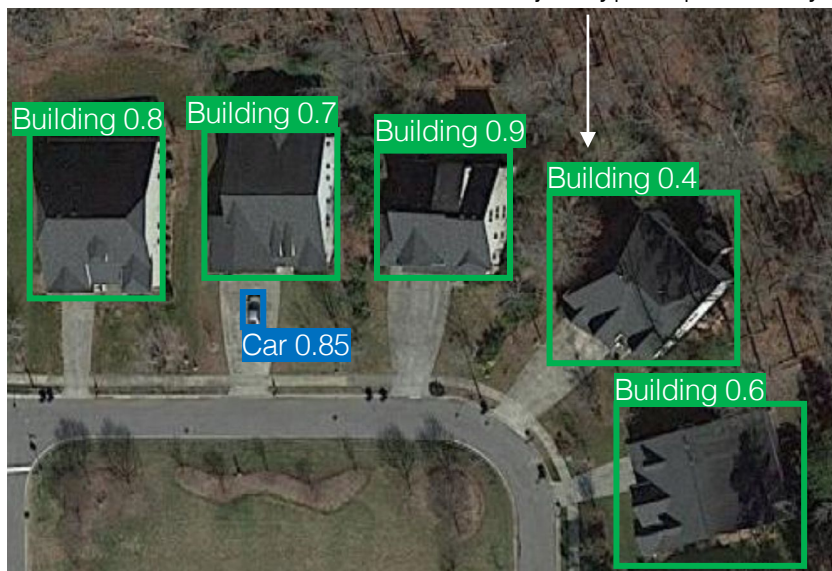


Image segmentation



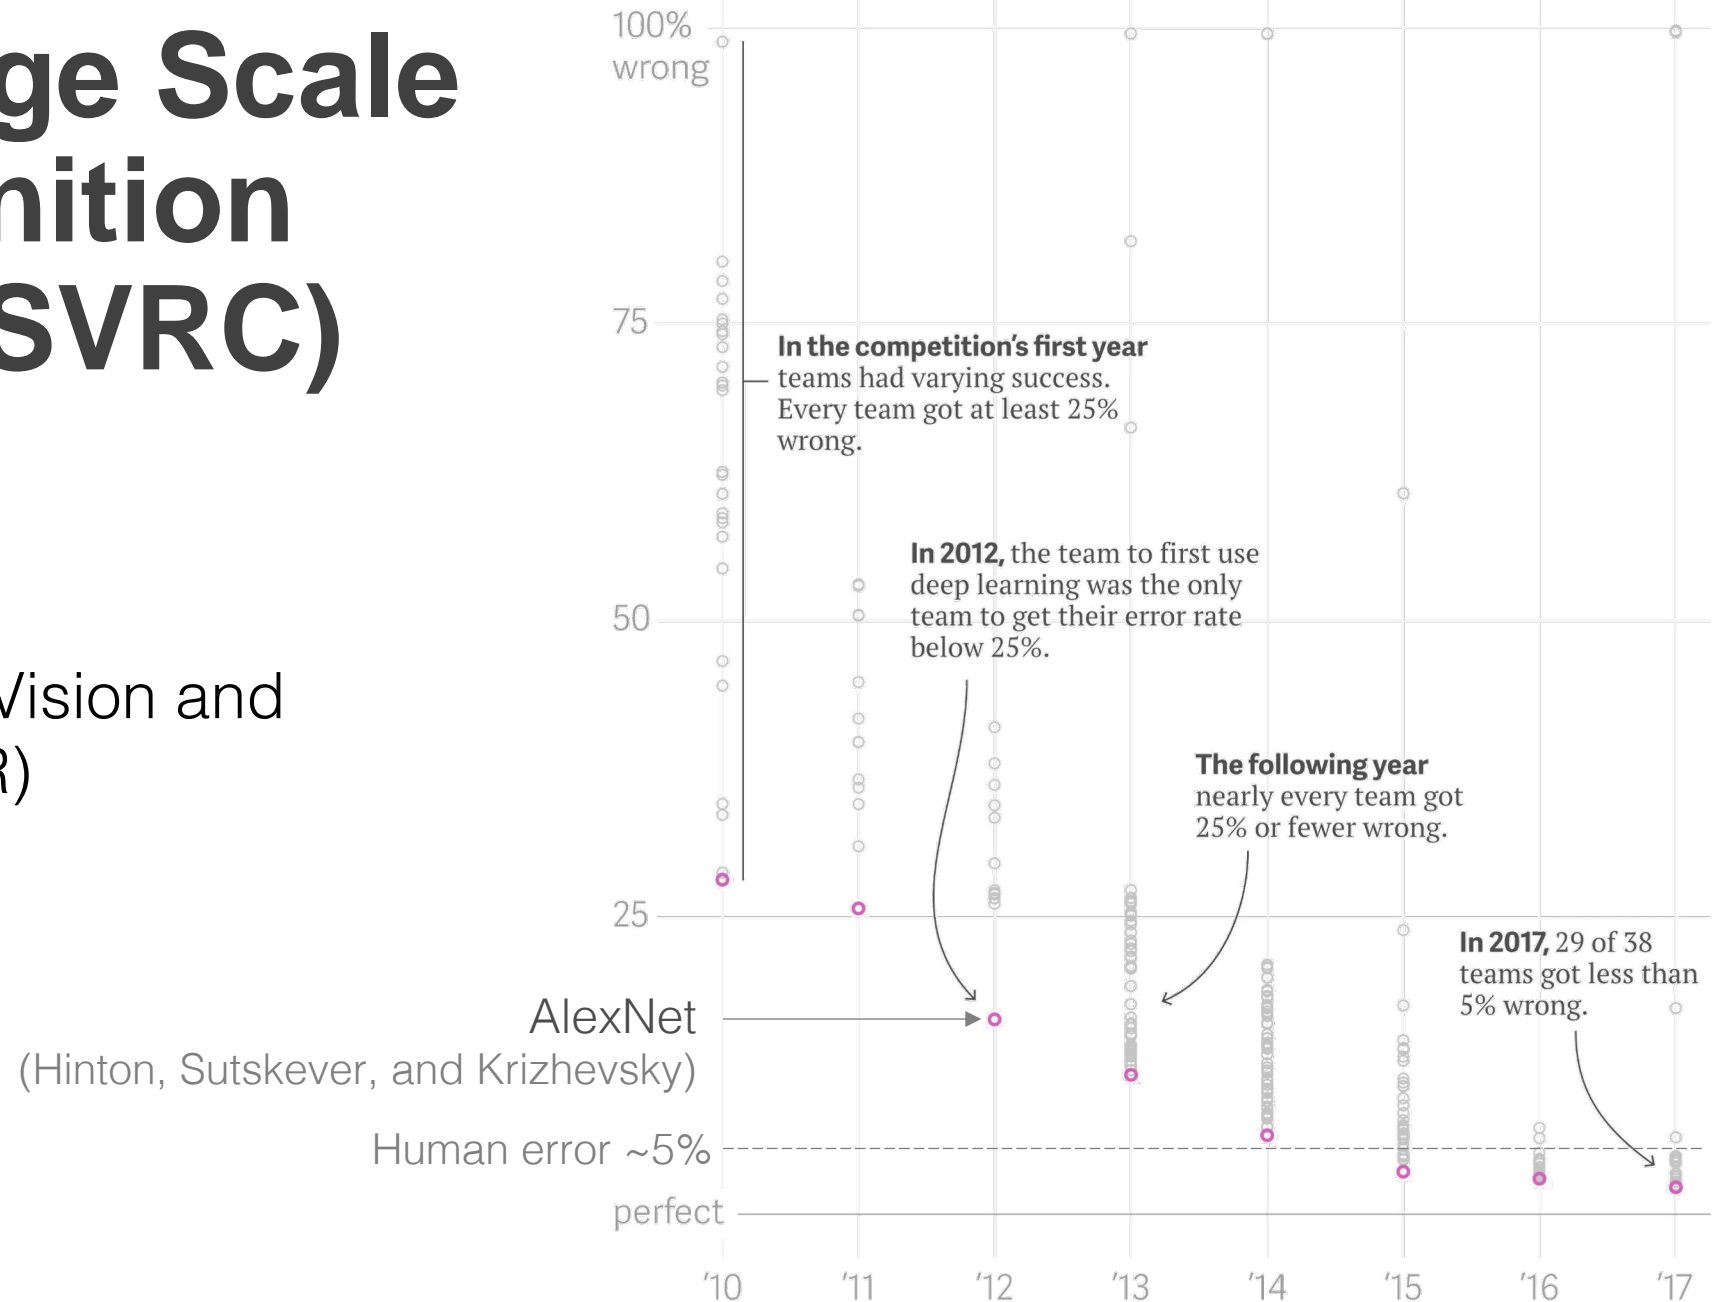
U-Net (2015)
SegNet (2016)
DeepLab (2017)
FCN (2016)



ImageNet Large Scale Visual Recognition Challenge (ILSVRC)

Fei-Fei Li et al. 2010 ([link](#))



Competition at:
Conference on Computer Vision and
Pattern Recognition (CVPR)



Source: Quartz, [link](#) David Yanofsky | Quartz

Data: ImageNet

Deep Learning Models Compared



Models compared for ImageNet


Many of these models are available through Keras ([link](#))

A. Canziani, E. Culurciello and A. Paszke, "Evaluation of neural network architectures for embedded systems," *2017 IEEE International Symposium on Circuits and Systems (ISCAS)*, Baltimore, MD, 2017, pp. 1-4.

Deep learning frameworks


Tensorflow ([link](#))

Framework for implementing graphical models, such as neural networks



Keras ([link](#))

Wrapper for Tensorflow to make coding easier: higher level and excellent API



Keras

PyTorch ([link](#))

Framework for implementing graphical models, such as neural networks



PyTorch

KERAS DEMO

Autoencoders

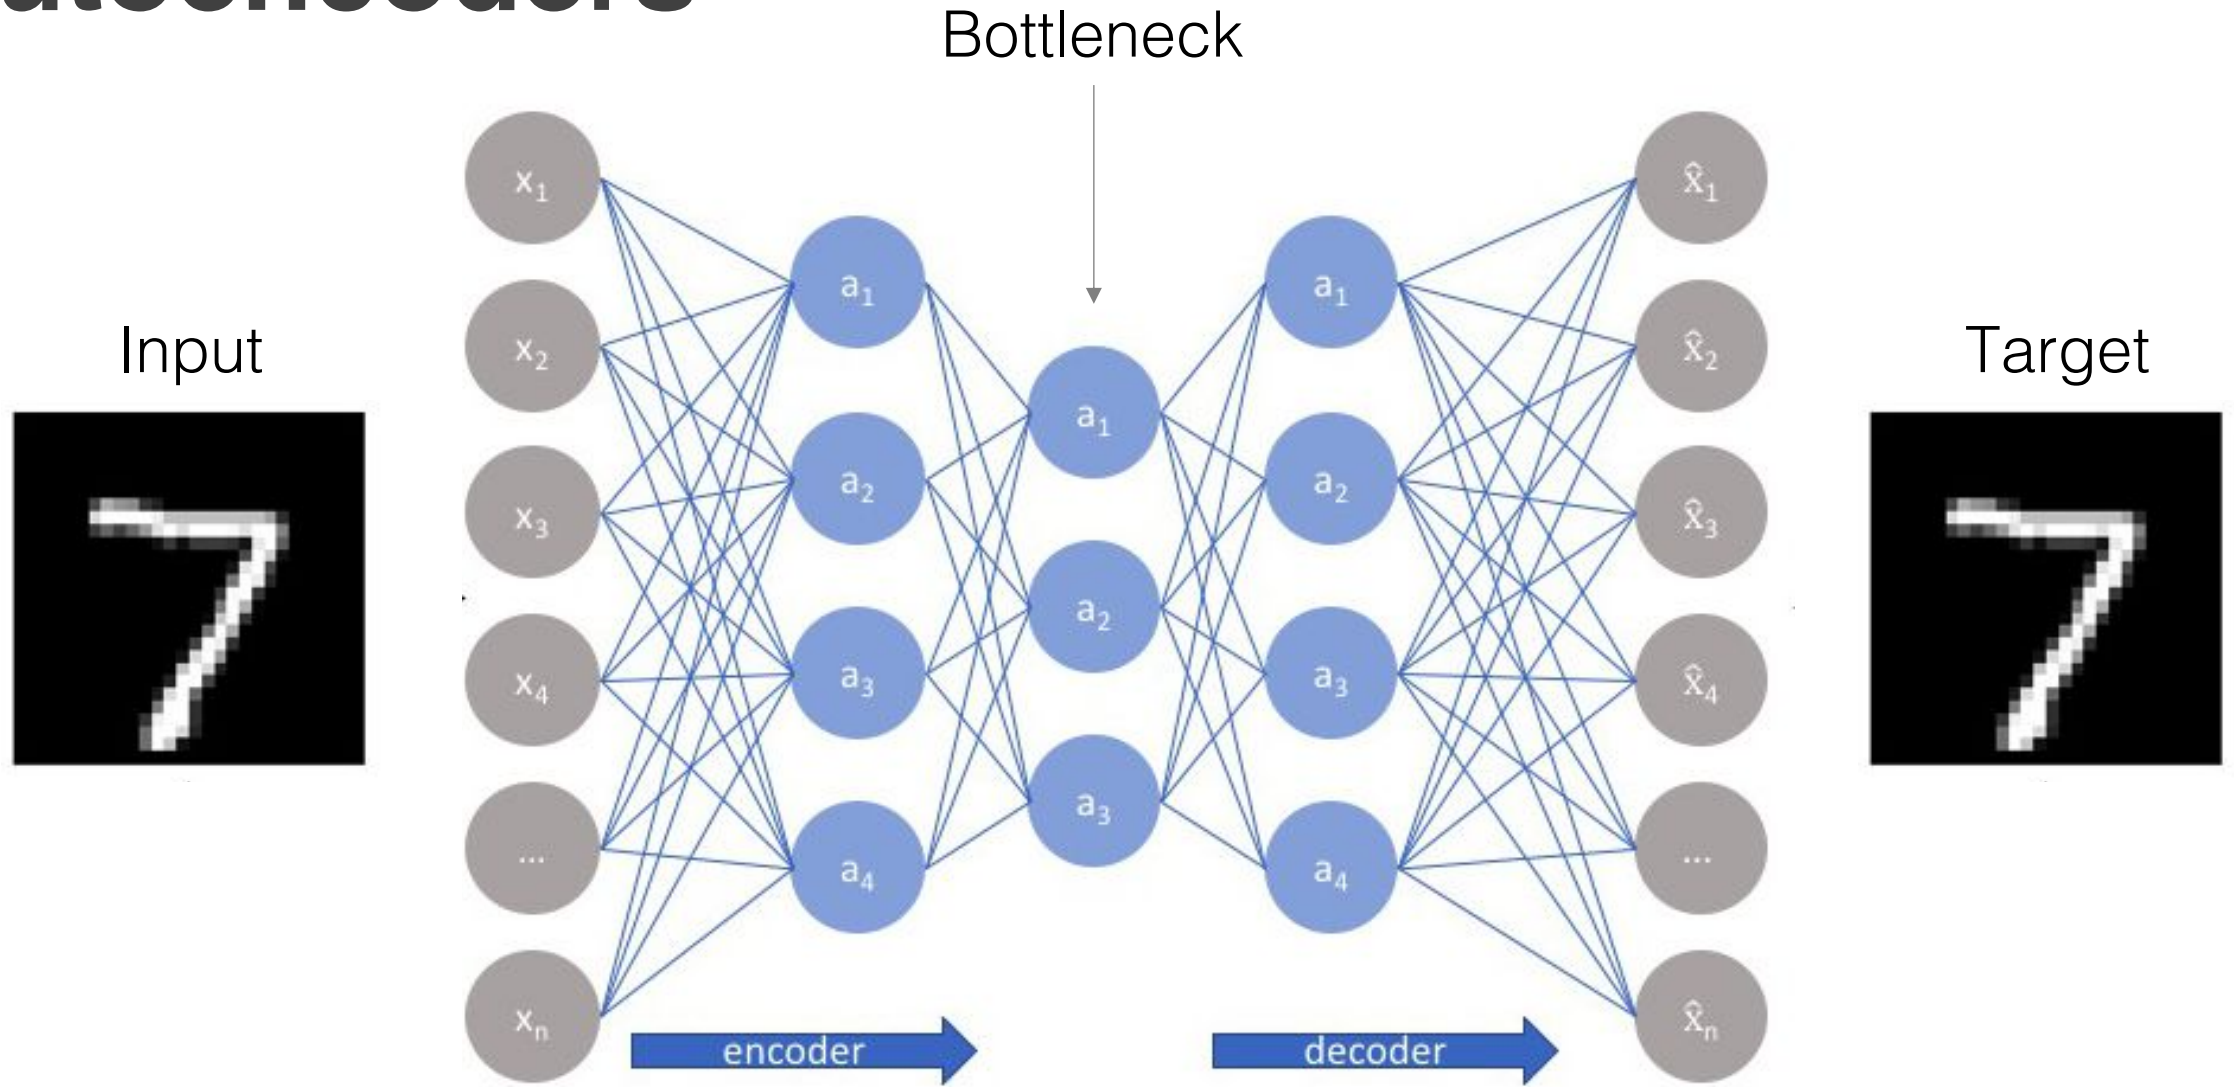



Image from: <https://www.jeremyjordan.me/autoencoders/>

Recurrent Neural Networks



x = inputs

o = outputs

h = hidden layers

U, V, W = model weights

Image from https://en.wikipedia.org/wiki/Recurrent_neural_network

Generative Adversarial Networks

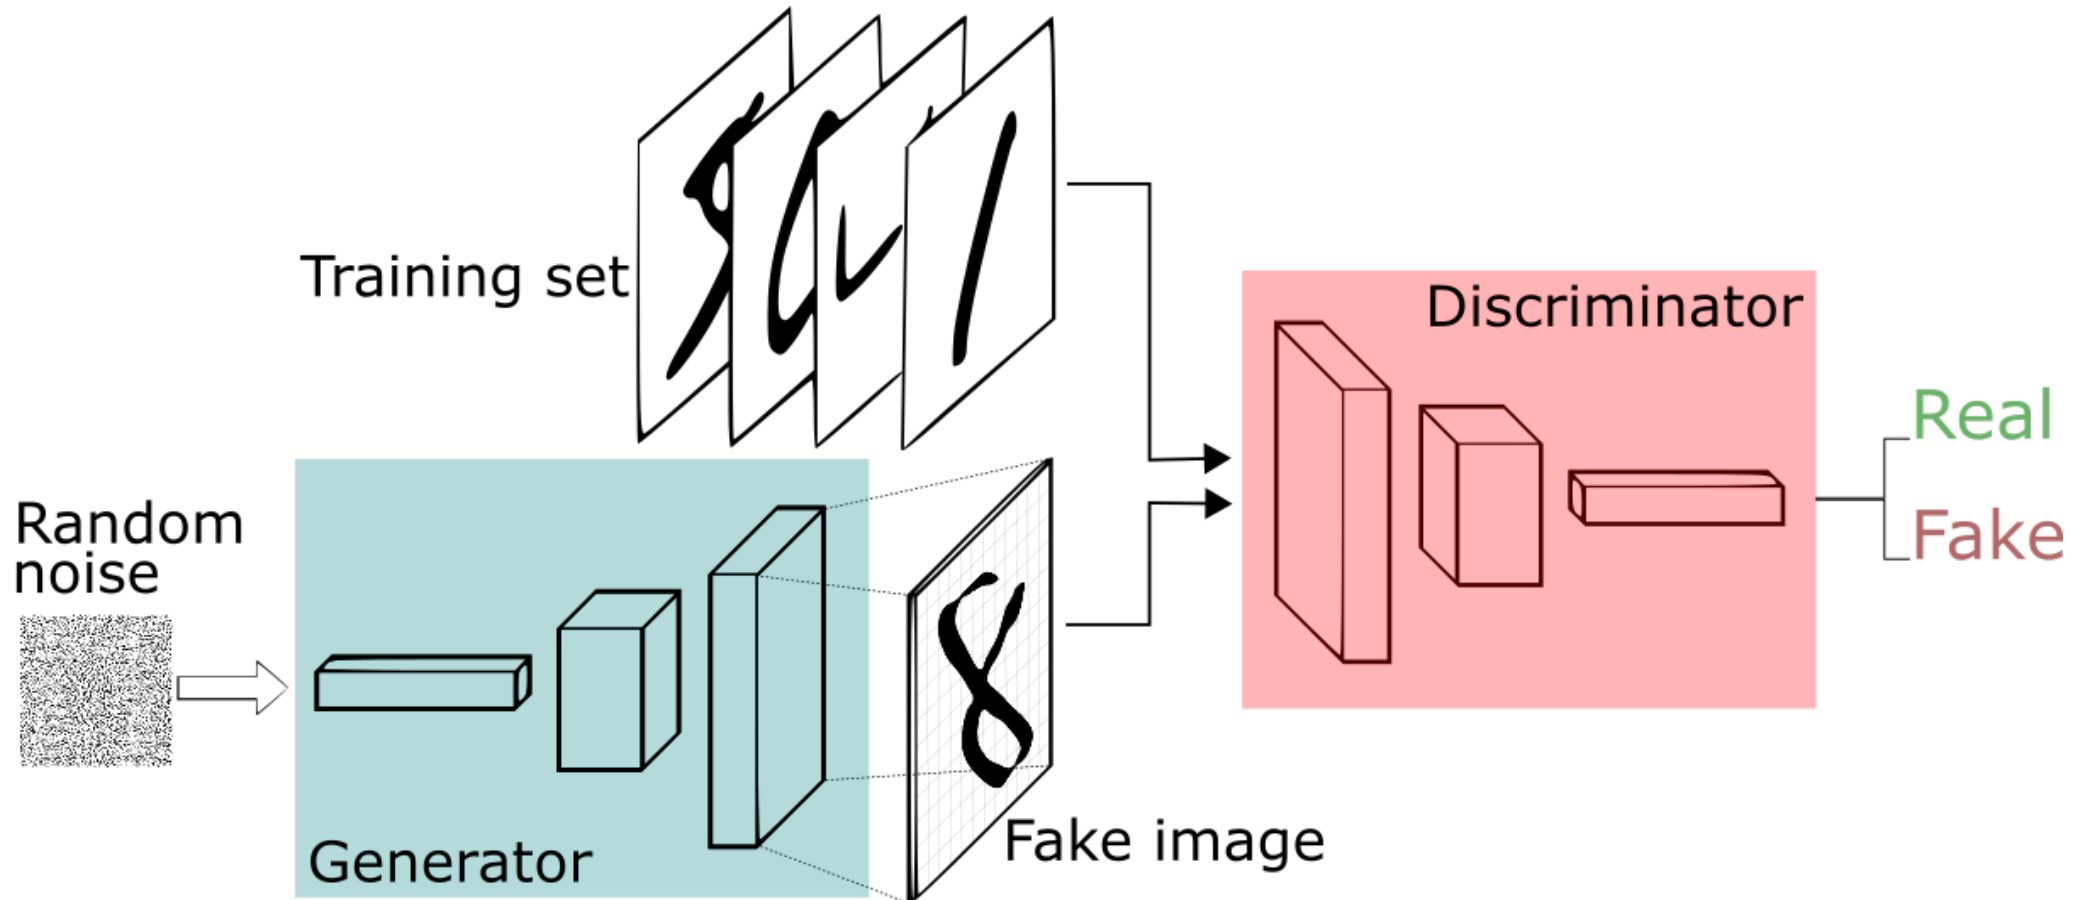


Image from: <https://skymind.ai/wiki/generative-adversarial-network-gan>

Supervised Learning Techniques

Covered so far

- Linear Regression
- K-Nearest Neighbors
- Perceptron
- Logistic Regression
- Linear Discriminant Analysis
- Quadratic Discriminant Analysis
- Naïve Bayes
- Support Vector Machines
- Decision Trees and Random Forests
- Ensemble methods (bagging, boosting, stacking)
- Neural Networks

Appropriate for:
● Classification
● Regression

Can be used with many machine learning techniques